Managed Agents x Valor

Action Plan: What to adopt, what to protect, what to test

April 2026

Governing Thought

Valor orchestrates. MA executes.

Adopt MA for the **hands**. Keep the **brain** local.

Hosted Redis is the bridge.

The Strategic Split

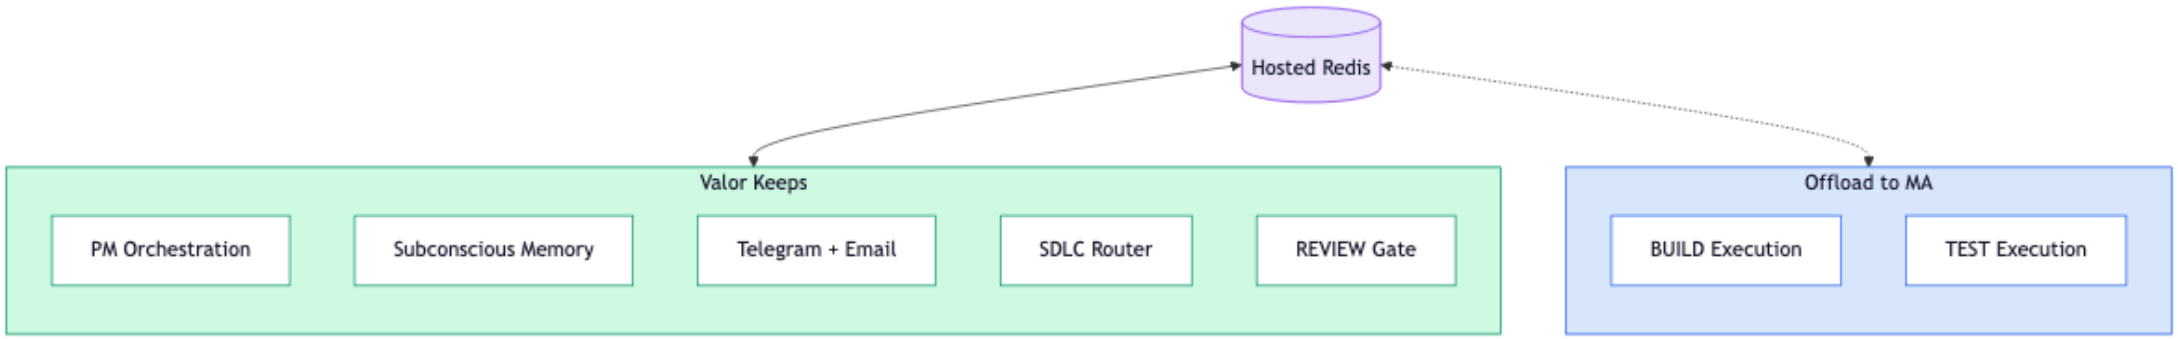

Valor keeps what MA cannot do: orchestration, memory, multi-channel delivery, the SDLC pipeline, and the final human-facing review gate.

MA replaces what Valor does poorly: sandboxed execution, environment management, crash recovery, scaling.

Why BUILD and TEST

	BUILD	TEST	EVERYTHING ELSE
Duration	Minutes to hours	Minutes	Seconds
Self-contained?	Takes plan in, branch out	Takes branch in, results out	Needs conversation context
Benefits from sandbox?	No local env pollution	Hermetic, no port conflicts	Not enough to matter
Sensitive?	No (code only)	No	MERGE/DEPLOY touch infra

BUILD is the clear winner. TEST follows naturally. The other 6 stages are fast, conversational, or security-sensitive.

Environments Persist, Sessions Don't

```
# Created ONCE — pre-baked with deps
environment = client.beta.environments.create(
    name="valor-fullstack",
    config={"type": "cloud", "networking": {"type": "unrestricted"}},
)

# Created PER TASK — starts fast, clean every time
session = client.beta.sessions.create(
    agent=build_agent.id,
    environment_id=environment.id,
    vault_ids=[github_vault.id, redis_vault.id],
)
```

Pre-bake Chromium, Playwright, Python deps into the environment. No rebuild per session.

Secrets injected via vault — never in context.

Headless Browser in Containers

MA containers can run headless Chromium. May be **more** reliable than local:

SCENARIO	CONTAINER	LOCAL
Backend unit/integration	Clean env	Works
Frontend unit (jsdom)	Clean env	Works
Frontend E2E (headless)	No port conflicts, no state leaks	Port collisions, background noise
Screenshot self-validation	Claude reads its own screenshots	Same
Mobile native	Not supported	Simulator only

The rule: MA runs tests + screenshots. Valor's REVIEW gate does the final visual check before merge. Never let a managed agent approve its own frontend work.

The Hard Problem

Subconscious memory doesn't port to Managed Agents.

Memory Portability Gap

Locally, the PostToolUse hook fires on every tool call. Bloom filter checks relevance. `<thought>` blocks inject into context. The agent never decides to recall — it just happens.

In MA, there are no hooks. No mid-generation injection. The API is request-response: once a turn starts, you can't insert tokens.

This is the **gating technical risk** for the hybrid architecture.

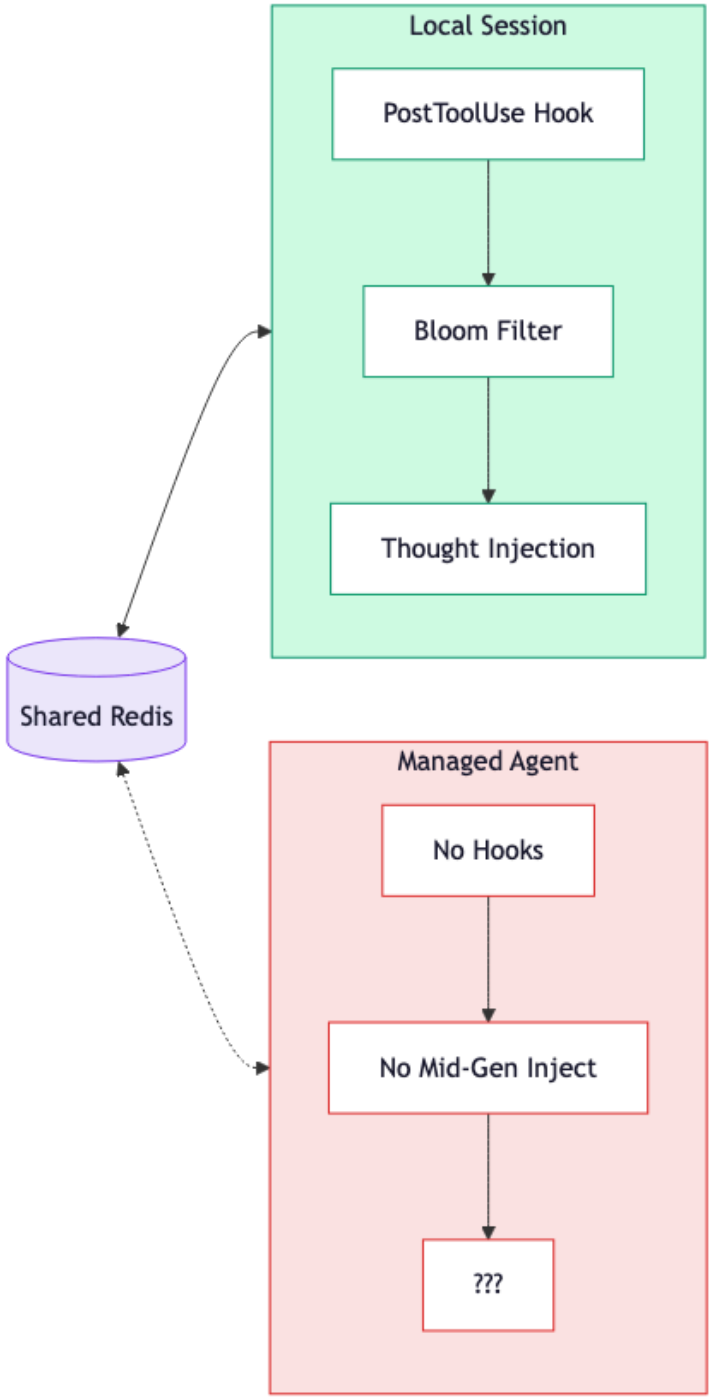

The Injection Spectrum

POINT	MECHANISM	SUBCONSCIOUS?	AVAILABLE IN MA?
Pre-Session	Memories in system prompt	Yes	Yes
Between Turns	`user.message` when idle	Yes	Yes (but only at turn boundaries)
Tool Result	Append `<thought>` to results	Yes	Only for custom tools
Mid-Generation	Inject during model output	Yes (local hooks)	**No**

We can't replicate the local experience exactly. But we can get close — if we control the right tool results.

Stealth Injection via Workflow Tools

The agent MUST call certain tools as part of its job: `run_tests`, `git_commit`, `submit_work`. These aren't memory tools — they're workflow tools that happen to route through Valor.

When the agent calls `run_tests`, Valor executes the tests, queries Redis for relevant memories, and returns test results **+** `<thought>` **blocks**. The agent sees verbose output. It doesn't know some of that output is injected memory. **Subconscious by design.**

> *Open question: Can `user.message` be sent during `running` state? If yes, we get true mid-turn steering and the workflow tool approach becomes a bonus, not the only path.*

The Session Lifecycle

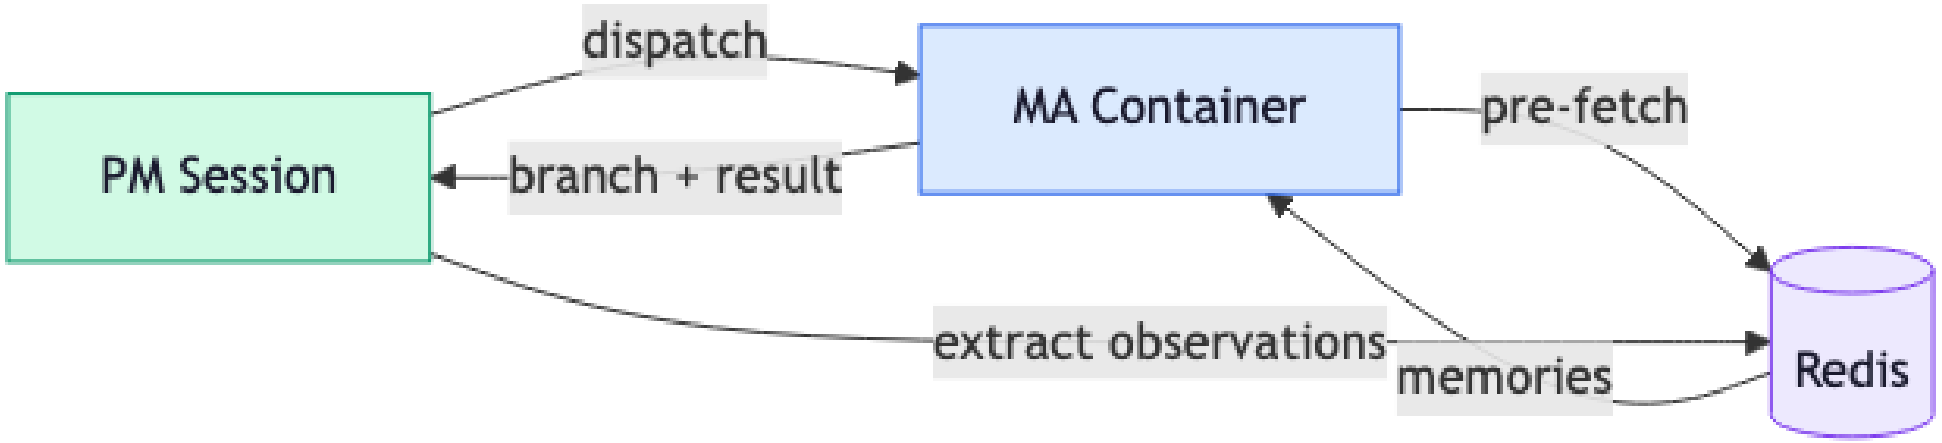

1. **PM pre-fetches** relevant memories into the MA session's first message
2. **MA executes** BUILD/TEST with access to hosted Redis
3. **Workflow tools** route through Valor for stealth memory injection
4. **Post-session**, Valor pulls the transcript and extracts new observations
5. **New memories** feed back to Redis — available to all future sessions

The memory system becomes a **shared brain** across both execution environments.

Quick Wins

1. Move to hosted Redis now

Not contingent on MA adoption. Unblocks multi-machine memory sharing today. Every improvement to Redis-based memory increases strategic value regardless of execution backend.

2. Stop investing in worktree edge cases

Worktree isolation, process management, crash recovery — all replaced by MA containers. Redirect effort to orchestration and memory quality.

> *Neither of these requires writing a single line of MA integration code. They're good decisions independent of the MA timeline.*

Experiment Pipeline

	EXPERIMENT	SUCCESS GATE	BLOCKS
A	Create MA environment with Chromium + Playwright + Python	Cold start < 30s	B
B	Run a real BUILD task, profile tool calls per turn	Avg < 5 calls/turn	C
C	Send `user.message` during `running` state	Mid-turn injection works	E
D	MA container connects to hosted Redis, verify memory search	< 50ms p95 latency	E
E	Full BUILD offload with pre-fetch + stealth injection + post-extraction	Quality matches local	Ship

A-B validate the platform. **C** is the critical fork — determines steering architecture. **D** validates the shared brain. **E** is end-to-end proof.

Sequential: each gates the next.

What to Watch

SIGNAL	WHAT IT MEANS FOR US	URGENCY
`user.message` during `running`	Unlocks true external steering — no need for workflow tool proxying	Test in Exp C
MA Memory (research preview)	Could replace our Redis memory system entirely	Watch closely
MA Multi-agent (research preview)	Could subsume PM/Dev session split	Medium
Hook-equivalent API	Makes subconscious memory fully portable	High — the real unlock
Pricing changes	Long BUILD sessions at $0.08/hr + tokens	Monitor

> ***Re-architecture trigger:*** *Memory stores GA + hook-equivalent API. Until both ship, Valor's memory is the moat.*

Summary

Now: Hosted Redis + stop worktree investment

Next: Experiments A through E, sequential

Watch: Hook API + Memory stores = migration trigger

The moat is orchestration and memory, not execution.