

OpenHuman vs. Hermes Agent

A side-by-side for agent builders

Two open agent stacks. Same era. Very different bets.

Why this comparison

Two new self-hosted agent stacks landed in the same window. Both claim "persistent memory," "tool use," and "autonomous skill." Both ship public docs and code.

But they are not competing for the same problem.

OpenHuman is a desktop app for one user's digital life. Hermes is a headless framework you deploy on a VPS and pipe through 20 chat platforms.

If you mix them up at the architecture stage, you waste a quarter.

At a glance

	OpenHuman	Hermes Agent
Made by	tinyhumans	Nous Research
Form factor	Desktop app (Tauri)	Headless framework
Primary surface	GUI + Obsidian vault	CLI + messaging gateway + ACP/IDE
Stack	Rust core + React UI	Python (run_agent.py , AIAgent)
Storage	SQLite + Markdown files	SQLite + FTS5
Distribution	One-user install, OAuth onboarding	Self-hosted, 6 execution backends
License	GNU GPL3	Open (Nous Research)

How each one thinks about "memory"

OpenHuman — Memory Tree

A deterministic ETL pipeline.

1. OAuth pulls from 118+ sources every 20 min
2. Normalize to provenance-tagged Markdown
3. Chunk into $\leq 3k$ -token segments
4. Score + entity-extract in background
5. Fold into per-source / per-topic / per-day summary trees
6. Browse in Obsidian. Edit by hand.

"No vector-soup black box."


Hermes — Curated + Dialectic

Two layers, both agent-driven.

1. **Agent-curated memory** — the agent decides what to persist, nudged periodically
2. **Honcho dialectic modeling** — separate service builds a user profile across sessions
3. **FTS5 recall** over session SQLite, with LLM summarization on retrieval
4. **Skills** — procedural memory, agentskills.io standard, self-improving


Memory is something the agent earns, not something piped in.

OpenHuman architecture



One process. One user. One machine (with cloud-backed inference).

Hermes architecture



Many entry points. Many backends. One `AIAgent` class.

Strengths

OpenHuman

- Transparent memory you can *read and edit*
- Onboarding is OAuth-easy
- Local-first; Ollama option for privacy
- TokenJuice keeps long-history costs sane
- Native desktop affordances (mascot, meeting agent, voice)
- Audit trail is just Markdown files

Hermes Agent

- Runs *headless* on \$5 VPS or GPU cluster
- Closed learning loop — skill creation + self-improvement
- 20+ chat platforms from one gateway
- MCP-compatible tool ecosystem
- Atropos hook for RL trajectory export
- Container isolation, command approval built in

Robustness — where the rubber meets the road

	OpenHuman	Hermes Agent
Failure domain	Single desktop process	Per-session, per-backend isolated
Sandboxing	Tauri process boundary; tools run in-process	Docker / Singularity / Modal containers
State recovery	SQLite + flat Markdown — trivially backupable	SQLite + FTS5 with lineage across compressions
Prompt stability	Implicit (single user, single session)	Explicit guarantee — system prompt frozen mid-conversation
Multi-user	No — one human, one install	Yes — per-platform isolation, user authorization
Blast radius	Your laptop	A container, a VPS, or nothing

OpenHuman's robustness story is "your filesystem is the backup." Hermes' is "everything runs in something you can kill."

Ideal use cases

Pick OpenHuman if

You are one knowledge worker, you want a second brain across Gmail / Slack / GitHub / Notion, and you want to see what the agent remembers.

Pick Hermes if

You are deploying an always-on agent that talks to many users on Telegram / Discord / Slack, runs jobs on remote infra, and you want skills to compound over time.

Pick neither if

You need a tightly-scoped coding agent inside an IDE (use Claude Code / Cursor) or a single API call wrapper (use the SDK directly).

The honest trade-off

OpenHuman bets on...

Transparency over autonomy.

The agent is bounded by what you let it see, and you can audit every chunk it reasons over. Skills are not really a concept — the system is the skill.

The cost: it scales to one user, on one machine, with one cloud subscription handling inference.

Hermes bets on...

Autonomy and reach over transparency.

Skills are first-class, self-improving, and shareable. Twenty chat platforms. Six execution backends. RL training hooks.

The cost: more moving parts, more services to operate, and "what does the agent know" is a query, not a folder.

Takeaways for an agent builder

1. **Memory shape follows form factor.** Desktop apps can afford a Markdown vault. Headless multi-tenant agents need indexed session stores with lineage.
2. **Sandboxing is a deployment choice, not a feature.** Hermes ships six backends because operators need that range. OpenHuman ships one because there is one operator.
3. **"Self-improving skills" is the live experiment.** Hermes is betting the agentskills.io ecosystem becomes a thing. Worth watching even if you don't adopt.

Questions?

OpenHuman: tinyhumans.gitbook.io/openhuman

Hermes: hermes-agent.nousresearch.com/docs