

The Skill Fleet Renovation

PR #1894: all 60 skills, two passes, eight parallel agents

Pass 1 fixed the structure. Pass 2 fixed the thinking.

The problem: skills rot silently

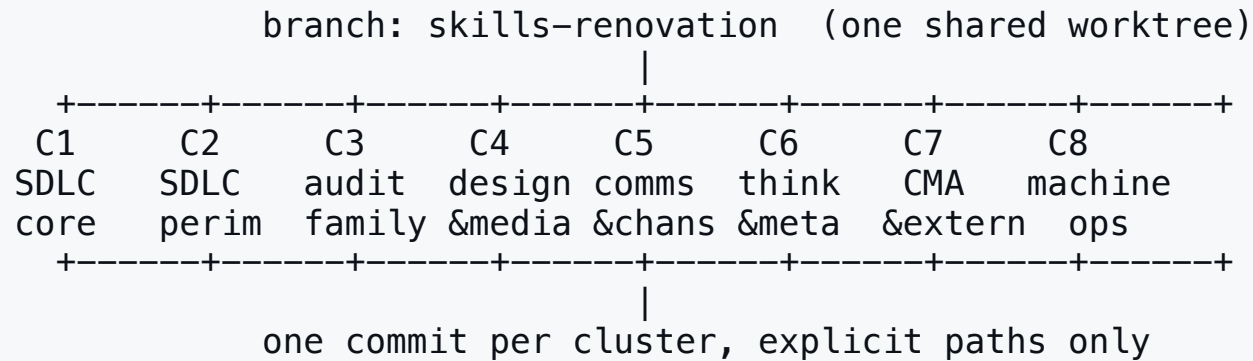
A skill is a markdown instruction file the agent loads on demand. Sixty of them had accumulated two years of drift, and nothing was watching.

- **7 monoliths**: SKILL.md bodies of 500 to 800 lines, fully loaded on every invocation
- **65 content warnings**: bloated descriptions, invalid frontmatter, junk files, orphaned sub-files
- **16,428 characters** of trigger descriptions injected into every system prompt
- **Live rot**: dead install scripts, a log format replaced long ago, pseudo-code APIs that never existed, hardcoded paths from another machine

Every stale line is a tax: wrong instructions get followed, and dead weight is paid for in tokens on every single call.

The method: fan out, own your lane

Both passes used the same shape: 8 cluster agents in one shared worktree, each owning a disjoint set of skill directories, judged by the five-lens rubric in `do-skills-audit/references/rubric.md`.



No two agents ever touched the same file, and the [skill-context seam](#) convention kept every global skill portable.

Two passes, two different lenses

	Pass 1: lint and structure	Pass 2: Opus-level optimization
Question asked	Is this skill <i>well-formed</i> ?	Does this skill <i>think well</i> ?
Targets	Monolith splits, description budgets, frontmatter, junk files, rot	Goal-first framing, over-instruction, duplicated rules, false claims
Typical fix	800-line body → 138 lines + sub-files	7-step micro-procedure → 4 goal-level steps
Held constant	Content (splits are reorganizations)	Determinism: commands, schemas, gates, probe sentences byte-exact
Verification	Audit lint, 76 unit tests	Line-by-line claim checks against live code and CLIs

Pass 2 assumes a smart reader. Hand-holding written for a weaker model is deleted; contracts and gates stay exact.

Scoreboard

Metric	Before	After
Skills audited	60	60
Line-count monolith FAILs	7	0
Content WARNs (rules 4 to 18)	65	0
Fleet description total	16,428 chars	10,463 chars
Remaining WARNs	72	52 (51 clear on next <code>/update</code> sync; 1 is the fleet description budget)

Verification: self-audit 15/15 PASS, 76 audit unit tests green, probe-sentence grep clean, 3-skill spot-check reads correctly from a Rust/JS repo.

The monolith splits (pass 1)

All content-preserving: bodies became a short SKILL.md plus load-on-demand sub-files with a "load when..." table.

Skill	Before	After	New structure
linkedin	800	138	4 references/ sub-files
x-com	728	101	3 references/ sub-files
do-pr-review	635	283	folded into sub-skills + outcome-contract.md
setup	620	153	5 phase-organized references/ sub-files
do-test	611	206	4 sub-files; orphaned PYTHON.md reconnected
do-design-system	534	295	3 references/ sub-files
do-build	526	186	deduplicated against existing sub-files
update (bonus)	338	143	modules.md + troubleshooting.md

Cluster 1

SDLC core

do-plan · do-plan-critique · do-build · do-patch · do-merge · do-sdlc · sdlc

do-plan

476 → 446 lines

- **Pass 1**: genericized the Telegram-bridge note and the xfail search so the skill reads correctly outside this repo
- **Pass 2**: fixed a real **Phase-3 numbering bug** (steps ran 1, 4, 5, 6), deleted a duplicate When-to-Use section, collapsed a 7-step WebSearch micro-procedure into 4 goal-level steps
- The commit-on-main rule was stated three times; now stated once

do-plan-critique

436 → 400 lines · description 222 → 128 chars

- **Pass 1:** description cut; flagged the in-body version history as per-invocation dead weight
- **Pass 2:** version history deleted, Step 2a's repo policy recitation moved to its `docs/sd1c/` seam file, critic cap rules deduplicated from 3 statements to 1
- **Found and fixed:** the seam file was stale, still describing the pre-#1714 six/seven-critic roster. Corrected to the real LITE(1) / FULL(3) split
- 3 barrier invariants from #1690 restored verbatim after the test reconciliation confirmed they were load-bearing

do-build

526 → 186 → 179 lines

- **Pass 1:** the biggest duplication find. Steps in SKILL.md repeated `WORKFLOW.md` and `PR_AND_CLEANUP.md` near-verbatim, so **every invocation paid roughly 2x tokens**
- Split resolved it: 186-line body, probe sentence kept verbatim
- **Pass 2:** deleted the Notes and Example-Execution walkthrough, removed "sent to Telegram chat" from the generic report template (bridge-specific wording in a global skill)

do-patch

351 → 310 lines · description 236 → 191 chars

- **Pass 1**: description trimmed while keeping all 5 trigger phrasings; test verification generalized across pytest, cargo, and npm
- **Pass 2**: deleted a duplicate Commit-Rules section, 24 lines of examples, and a Context-Awareness section
- **Restored** the test-anchored heading "Sync Plan Checkbox...", which fixed 2 of 3 stale content-test failures on the spot

do-merge

178 lines · description 332 → 155 chars

- **Pass 1:** description halved
- **Pass 2:** removed the last repo incantation (`touch data/merge_authorized_{PR}`) from the generic body; it now lives only in the seam file
- Cluster verdict: **reference-quality skill**. The verify-then-merge gate needed no thinking repairs
- Disposition note: Steps 1 to 3 are a plausible future `scripts/` extraction, but UNKNOWN-retry and dependabot branching still carry judgment

do-sdlc

153 lines · description 340 → 176 chars

- **Pass 1:** description halved; surfaced a real governance discrepancy
- **Pass 2:** every claim verified against current code. Nothing to cut

Discrepancy for #1883: `docs/features/skill-context-convention.md` lists do-sdlc as Bucket C project-only, but it lives in `skills-global/` and syncs to every machine. Either the doc is stale or the #1783 move was never executed. Needs human sign-off.

sdlc (project router)

263 → 248 lines

- **Pass 1:** no changes; flagged the G1-G7 guard table as duplicating `sdlc-tool next-skill`
- **Pass 2:** verified **#1558 is closed** with `ensure=True` on all write paths, so the Step 1.5 belt-and-suspenders `session-ensure` and its `REDUNDANT-AFTER-#1558` marker were retired
- Guard table kept byte-exact (a parity test asserts it) but reframed as an interpretation reference: the tool is the sole dispatch authority
- Command discipline rule deduplicated from 3 statements to 1

Cluster 2

SDLC periphery

do-test · do-pr-review · do-docs · do-issue · do-investigation-issue

do-test

611 → 206 → 201 lines

- **Pass 1**: monolith split into 4 sub-files (parallel-dispatch, baseline-verification, quality-gates, special-targets); the orphaned `PYTHON.md` is finally referenced
- **Pass 2**: the repo's source-to-test mappings, generalized out of the body in pass 1, were **restored into** `docs/sdlc/do-test.md` so the knowledge survives at the seam
- Fixed a stale claim in `PYTHON.md`: lint default is ruff, the doc still said black
- Constants and quality gates verified load-bearing and kept exact

do-pr-review

635 → 283 → 248 lines

- **Pass 1:** content folded into existing sub-skills plus a new `outcome-contract.md`; junk `sub-skills/README.md` deleted; argument-hint added
- **Pass 2:** the **last two repo couplings pulled through the seam** (`plan_checkbox_writer` command and the `requires_real_chrome` flag now live in `docs/sdlc/do-pr-review.md`)
- Fixed a stale reference to `.claude/commands/do-merge.md`; the no-blank-verdicts rule went from 3 statements to 1
- All 8 Hard Rules kept word-for-word

do-docs

307 → 305 lines · description 236 → 139 chars

- **Pass 1**: description cut by 40 percent
- **Pass 2**: mostly a verification pass. Every seam automation hook was checked against live code and confirmed unaffected by the #1828 reflection-scheduler split that merged mid-branch
- One doubly-stated merge rule deduplicated

The quiet skills matter too: confirming a skill is *accurate* is the same work as fixing one that is not.

do-issue

174 → 167 lines · description 256 → 142 chars

- **Pass 1:** description cut; the "track this" trigger promoted from the body into the description where it can actually fire
- **Pass 2:** restored a lost **### Step 4 heading**. RECON.md pointed readers at a step that no longer existed in the body
- Goal statement moved above the probe sentence; duplicate When-to-Use deleted

do-investigation-issue

133 → 110 lines · description 429 → 193 chars

- **Pass 1**: description cut by more than half; overlap with do-issue examined and deliberately kept separate, since a merge would blunt trigger precision
- **Pass 2**: fixed a **live template-drift bug**. The inline heredoc copy of the issue template had silently diverged from `TEMPLATE.md`
- Now `TEMPLATE.md` is the single source of truth; the body references it instead of embedding a copy

Cluster 3

Audit family

do-skills-audit · audit-hooks · audit-models · audit-tools · do-integration-audit · do-oop-audit · new-audit-skill

do-skills-audit

92 lines · untouched in both passes

- The auditor that defined this renovation's rubric was placed under a **light-touch mandate**: you do not rewrite the measuring stick mid-measurement
- Self-audit exits 0 with 15/15 checks passing, before and after
- Its rubric (`references/rubric.md`) drove all 8 clusters in both passes

One finding *about* it: rule 13's coupling-signal set misses `python tools/...` invocations, which is how do-discover-paths escaped detection. Extending the signal set is a #1883 disposition item.

audit-hooks

105 → 98 lines + new 24-line seam file

- **Pass 1**: description 275 → 156 chars
- **Pass 2**: rewritten goal-first around its real invariant: *no hook may hang a session*
- **Rot fix**: the body listed 7 validators; the repo has **21**. The stale list was replaced with a live directory enumeration declared in a new `.claude/skill-context/audit-hooks.md`
- Now the skill can never drift from the validator roster again

audit-models

102 → 92 lines + new 28-line seam file

- **Pass 1:** added the probe sentence and reframed hardcoded `project_key` / `ReflectionIgnore` specifics as overridable defaults
- **Pass 2:** the seam file was **actually created** (pass 1 had added only the probe). It declares Popoto, the `project_key` `KeyField`, `ReflectionIgnore`, and the no-raw-Redis rule; the body got honest generic defaults

Merge candidate for #1883: heavy check overlap with do-oop-audit. Either merge in, or move project-only. Both need a `RENAMED_REMOVALS` entry.

audit-tools

96 → 81 lines · description 320 → 179 chars

- **Pass 1**: description cut by 44 percent
- **Pass 2**: intro merged goal-first; version-history section deleted
- Trigger collision flagged: "reviewing tool quality" fires both this skill and the user-level `audit-next-tool` orphan.
Disposition pass must pick one owner

do-integration-audit

291 → 208 lines

- **Pass 1:** missing trigger phrase added (rule 4)
- **Pass 2: 83 lines of Django and FastAPI worked examples distilled into a 10-line, 4-finding style reference.** An Opus-class reader needs the shape of a good finding, not three full essays
- The verify-before-report discipline kept word-for-word: it is the skill's contract

do-oop-audit

218 → 158 lines · description 422 → 193 chars

- **Pass 1:** description cut by more than half; the "audit models" trigger deliberately dropped because it collided with the audit-models skill
- **Pass 2:** 55 lines of worked examples became a 5-finding style reference
- All 14 checks kept exact
- Designated the survivor if audit-models merges in under #1883

new-audit-skill

136 → 128 lines · description 439 → 194 chars

- **Pass 1:** `scripts/audit.py` looked like a broken reference but is a template placeholder, so it is now placeholder-fenced; "8 existing audit skills in this repo" genericized
- **Pass 2: reference-table rot fixed:** do-skills-audit grew from 12 to 20 rules, audit-tools became a 10-check report-only skill, and a duplicate docs-auditor row was merged
- Disposition: candidate to fold into new-skill as `references/audit-template.md`

Cluster 4

Design and media

*do-design-system · do-design-audit · frontend-design · pencil-design · do-presentation · mermaid-render ·
do-debrief · do-voice-recording*

do-design-system

534 → 295 → 261 lines · description 482 → 198 chars

- **Pass 1:** monolith split into 3 references/ sub-files (file-organization, moodboard-capture, pen-editing) with a load-when table covering all four sub-files
- **Pass 2:** goal-first rewrite with explicit success criteria: 3 to 7 charter-grounded additive edits per pass
- A duplicate charter-HALT paragraph deduplicated; all safety gates verbatim

do-design-audit

258 → 229 lines · description 348 → 198 chars

- **Pass 1**: description trimmed
- **Pass 2**: rewritten deliverable-first; duplicate When-to-Use and Rating-Scale sections deleted
- The 10-dimension rubric contract untouched

Near-duplicate pair: do-design-audit and the user-level orphan do-design-review are almost the same skill. #1883 must pick a survivor.

frontend-design

132 → 63 lines · description 253 → 150 chars

- **Pass 2's philosophical fix**: the body tried to teach *taste by enumeration*, roughly 45 DO/DON'T bullets
- Replaced with **design intent plus sharp anti-goals per section**: say what good looks like and what to refuse, and trust the model's judgment in between
- Every AI-slop fingerprint that do-design-audit cross-references was preserved
- The 7 reference files were already right and went untouched

pencil-design

196 → 194 lines · description 470 → 181 chars

- **Pass 1:** description cut 61 percent, keeping the NOT-trigger boundary that stops it firing on ordinary CSS work
- Removed the repo-specific "Distribution" section: hardlink-sync mechanics are always-true policy already covered by CLAUDE.md, the first `flagged_for_claude_md` finding of the fleet
- **Pass 2:** doubled preamble folded; all failure-history rails kept exact

do-presentation

340 → 323 lines · description 319 → 171 chars

- **Pass 1**: description trimmed; brand-logo step flagged as the first sub-file candidate if the body grows
- **Pass 2**: goal-first opening now tied directly to the Step-10 verify checklist; version history deleted; Marp and seam contracts kept exact
- Fixed a **phantom cross-reference** in CONTENT_GUIDE.md pointing at CSS classes that no longer exist

This deck was produced by the renovated skill.

The quiet ones: mermaid-render · do-debrief · do-voice-recording

Skill	Pass 1	Pass 2
mermaid-render	description 324 → 179; byob doc link verified live	unchanged: already renovation-grade
do-debrief	description 281 → 175; seam indirection untouched	anti-patterns 7 → 3; the rest live once in their owning phases
do-voice-recording	description 392 → 189; a 38-line model citizen	unchanged: reference implementation of the seam pattern, seam verified accurate

The fleet's target state looks like do-voice-recording: tiny, portable, seam-aware, and true.

Cluster 5

Comms and channels

*email · google-workspace · telegram · checking-system-logs · reading-sms-messages · linkedin · x-com ·
authenticity-pass · sentry*

linkedin

800 → 138 → 131 lines

- **Pass 1:** the fleet's largest monolith, split into 4 references/ sub-files (dom-model, messages, posting, feed-engagement); diff-verified content-preserving
- **Pass 2:** duplicate BYOB stack explanation collapsed; Notes deduplicated; all DOM selectors exact
- The em-dash prohibition stays inline even though it duplicates user-global policy, because drafter subagent prompts need it verbatim
- Disposition: linkedin and x-com share a large skeleton (drafter delegation, persona cold-read loop, authenticity gate)

X-com

728 → 101 → 97 lines

- **Pass 1**: second-largest monolith, split into 3 references/ sub-files (dms, posting, timeline-engagement)
- **Pass 2**: two near-duplicate "Iterate: 3-4 rounds" sections inherited from the original were **consolidated into one 10-step loop** (posting.md 286 → 268)
- Shared-skeleton extraction with linkedin was evaluated and **deliberately rejected**: 5-vs-4 personas and inverted hashtag rules mean the identical residue is about 30 lines, less than the indirection would cost

google-workspace

290 → 107 lines

- **Pass 1:** `user-invocable: false` added, since the skill exists for the model, not the slash menu
- **Pass 2:** the **largest single cut in the fleet**. The body read like a human-facing behavior guide; it is now contracts and recipes
- Kept exact: the tool-selection ladder, draft-first email policy, and composition constraints

Deleting 63 percent of a skill while keeping 100 percent of its contract is the pass-2 thesis in one slide.

telegram

223 → 217 lines

- **Pass 1**: already PASS, no changes
- **Pass 2**: every documented flag **empirically verified against the live CLI**
- Added the missing `--reply-to`, `--voice-note`, and `--await-reply` documentation
- Corrected a stale claim that sends go "through Telethon directly": they route through the bridge relay now

checking-system-logs

66 → 79 lines: it grew, on purpose

- **Pass 2's clearest rot find:** the skill was **fully rotted**. Every recipe grepped `bridge.events.jsonl`, a file replaced long ago by the BridgeEvent Redis model
- An agent following this skill would have searched a file that no longer receives events
- Rebuilt from live surfaces: `analyze_logs.py`, Popoto queries, the current log table including `reflection_worker.log`, and `valor_session telemetry`
- One of two skills that got *longer* in pass 2. Accuracy beats brevity

sentry

92 → 70 lines

- **Pass 1 rot fix:** hardcoded `/Users/valorengels/src/ai` replaced with `~/src/ai`, so the skill works on machines with other usernames; argument-hint `[--apply]` added
- **Pass 2:** two 20-line `python -c` blocks that differed only by `SENTRY_TRIAGE_APPLY=1` were **consolidated to one**, verified against the real gating in `reflections/sentry_triage.py`
- Disposition: the diverged user-level sentry copy and the `sentry-cli` orphan still need #1883 resolution

email · reading-sms-messages

Skill	Pass 1	Pass 2
email	description 268 → 196, all triggers kept, probe untouched	70 → 67; the BYOB prohibition was stated 3 times, now once; seam verified current against <code>tools/valor_email.py</code>
reading-sms-messages	already PASS, no changes	every command verified against the live surface; untouched

Verification without change is still a deliverable: these two are now *known* accurate, not assumed accurate.

authenticity-pass

137 → 129 lines · description 274 → 173 chars

- **Pass 1**: description rewritten with a proper "Use when"; argument-hint `[draft-file-path]` added
- **Pass 2**: rationale prose compressed to one cited paragraph
- The PASS/BLOCK contract kept **byte-exact**: linkedin and x-com gate on it before anything goes live
- Model-tier proposal for #1883: this is a classification task, sonnet would suffice

Cluster 6

Thinking and meta

*analyze · grill-me · zoom-out · ontologies · weekly-review · reclassify · skillify · new-skill · pthread · tdd ·
deepen · observability · claude-standards*

pthread: the fleet's biggest rewrite

221-char description → 138 lines → 26 lines

- **Pass 1** trimmed the description and flagged the real problem
- **Pass 2** confirmed it: the body was **pseudo-code fiction**. `spawn_subagent`, `thread_metrics`, and a wrong `run_in_background` claim describe an API that does not exist
- Full rewrite against real Agent-tool mechanics: batch independent calls in one message, self-contained prompts, never parallelize dependent work, worktree isolation for concurrent writers

flagged_for_claude_md: CLAUDE.md Development Principle 8 duplicates this skill's core rule. The policy lives in two places.

analyze

35 → 43 lines · description 556 → 192 chars

- **Pass 1:** the fleet's most bloated description cut by two thirds; overflow triggers and the do-not-trigger guard moved verbatim into the body opening
- **Pass 2: graceful degradation added.** The skill depends on the `strategic-analyst` agent that only exists in this repo's `.claude/agents/` ; on other machines it now runs the identical 5-lens protocol inline via general-purpose agents
- The second skill that grew in pass 2, and for the same reason: correctness

skillify + new-skill: template convergence

skillify 142 → 91 · new-skill 118 → 113 + 56-line WORKFLOW_TEMPLATE.md

- **Pass 1:** skillify carried invalid frontmatter (`when_to_use` , `arguments`), migrated to `description` / `argument-hint` ; its embedded generated-skill template updated to valid fields
- **Pass 2:** skillify's **drifted 54-line embedded template deleted**; it now defers to new-skill via a relative link, and both sync together
- new-skill is now the canonical owner of both skill shapes (`SKILL_TEMPLATE.md` + new `WORKFLOW_TEMPLATE.md`)
- Bonus: the version-history prescription was removed from the template, so newly generated skills stop shipping changelogs

reclassify

52 → 32 lines + new 24-line seam file · description 219 → 164 chars

- **Pass 1:** flagged that the skill assumes this repo's `docs/plans/*.md` frontmatter convention while living in `skills-global/`
- **Pass 2:** resolved it with the seam pattern. A probe sentence was added, and `.claude/skill-context/reclassify.md` now declares the enforced type values, the immutability lock, and the validator hooks
- The body shrank 38 percent because the repo-specific half moved to where it belongs

weekly-review

138 → 104 lines · description 292 → 197 chars

- **Pass 1**: description trimmed; minor flags noted (saves to `/tmp/`, offers macOS TextEdit)
- **Pass 2**: a 12-item category menu replaced with **a principle plus 5 examples**. Enumerating every possible category is over-instruction; stating the categorization principle transfers better
- Version history dropped

claude-standards

123 lines · description 359 → 194 chars

- **Pass 1**: description cut by 46 percent
- **Pass 2**: two repo-parochial inventory rows made portable
- Disposition: it overlaps do-skills-audit on the skills domain. Proposal is to carve that domain out and delegate, and its 7 references/ guides arguably form a separate background reference skill

Already sharp: six skills, near-zero deltas

Skill	What happened
grill-me	desc 230 → 188 in pass 1; zero pass-2 changes
tdd	repo-agnostic lint/format generalization only; zero pass-2 changes
observability	desc 266 → 197; zero pass-2 changes
zoom-out	desc 235 → 189; one duplicated trigger sentence removed
deepen	desc 245 → 176; read-only rule deduped 3x → 1x
ontologies	desc 223 → 182; repo-domain example swapped for a generic Order checkout example

A renovation that changes nothing where nothing is wrong is evidence the rubric measures the right things.

Cluster 7

CMA and external

imagine-agent · build-agent · do-discover-paths · computer-use · officecli · ebook-ingest

do-discover-paths

213 → 106 lines + new 56-line seam file

- **Pass 1** found the governance gap: the body invokes this repo's toolchain (`tools/happy_path_*.py` , Rodney scripts) while living in `skills-global/` , and rule 13 missed it because `python tools/...` is outside the coupling-signal set
- **Pass 2** applied the seam convention: probe added, `.claude/skill-context/do-discover-paths.md` declares the validator, generator, Rodney, and `BYOB_ALLOW_EVAL` specifics
- Generic baseline on foreign machines: write the trace and stop
- Trace schema, actions, and selector rules kept exact; a duplicated minified JS one-liner deleted

officecli

412 → 420 lines: grew via verification

- **Pass 1:** description trimmed to 200 chars; argument-hint added
- **Pass 2: every documented command empirically run against the live CLI v1.0.128.** The installed 1.0.29 binary rejects half the documented surface
- Added a stale-binary guard: on "Unrecognized command", re-run the installer
- Documented three commands that existed but were absent from the skill (`save` / `import` / `merge` , plus `svg/screenshot/forms` views and the resident-mode flush footgun)

imagine-agent

103 → 101 lines · description 495 → 193 chars

- **Pass 1 rot fix:** `references/build-sheet.md` was a **mis-written relative path** that resolved nowhere. Corrected to `../build-agent/references/build-sheet.md`, which works both in-repo and in synced installs
- The imagine-agent → build-agent handoff contract stayed intact
- **Pass 2:** duplicated opening merged; the interview protocol confirmed judgment-first

ebook-ingest

762-char description → 198 · body 444 → 289 lines

- **Pass 1**: the fleet's longest description cut 74 percent, NOT-triggers moved to the body opening
- The interesting fix: two prose blocks saying "save this as `scripts/X.py`" were **extracted into real bundled scripts** (`annas_get.py` , `clean_book.py`), compile- and ruff-verified
- Code that ships as prose is code that silently rots; now it is code
- **Pass 2**: overview rewritten deliverable-first around `library/processed/<slug>.md` with success criteria

Untouched by design: build-agent · computer-use

Skill	Pass 1	Pass 2 verdict
build-agent	description 369 → 194	untouched: its API references are battle-scarred contracts , every line earned by a real failure
computer-use	description 310 → 193; probe verbatim	untouched: model citizen of the seam convention , all seam claims verified

- build-agent carries a #1883 model-tier proposal: fable, because a client-facing launch has billing consequences

Cluster 8

Machine ops

setup · update · prime · do-deploy · do-deploy-example

setup

620 → 153 → 154 lines

- **Pass 1:** monolith split into 5 phase-organized `references/` sub-files; dead `install_reflections.sh` / `reflections.py` references replaced with the then-current worker scheduler reality
- **Pass 2:** that pass-1 fix was **already invalidated**. #1828 merged mid-branch and split reflections into their own subprocess
- Step 8 now installs all three services (worker, reflection-worker, sdlc-reflection)

A fresh machine set up from the pass-1 text would never have run reflections. Renovation is not a one-time event.

update

338 → 143 → 120 lines

- **Pass 1:** bonus monolith split (`modules.md` + `troubleshooting.md`); `disable-model-invocation: true` added; same reflections rot fixed as setup
- **Pass 2:** reconciled with #1828: installs and verifies `com.valor.reflection-worker`, probes with `python -m reflections --dry-run`, watches `reflection_worker.log`
- Stale unbuilt BYOB bump/rollback machinery in `modules.md` replaced with a design pointer
- New troubleshooting section for the reflection scheduler

prime

104 → 94 lines

- **Pass 1**: no changes, one stale-comment flag
- **Pass 2**: the onboarding guide was **describing a two-year-old architecture**: no Redis queue, no worker, no reflection subprocess, and an orphaned permission-model table
- Rebuilt around the current flow: bridge → Redis queue → worker → SDK / granite-PTY, plus out-of-process reflections
- The C-Thread/L-Thread taxonomy deleted; an agent-integration rule added (`[project.scripts]` visibility)

do-deploy · do-deploy-example

Skill	Pass 1	Pass 2
do-deploy	description 231 → 146	140 → 142: Step 3 broadened to bridge and worker health; an overclaiming comment fixed; report template gains a worker-status line
do-deploy-example	trigger phrase added; wrong template copy source fixed (skills/ → skills-global/)	unchanged: every line is contract, safety rail, or customization guide

Aftermath

What broke, what got flagged, what it taught us

The test reconciliation

Generalizing skill bodies broke **30 unit tests** that grep skill prose. A dedicated repair agent classified every failure instead of patching blindly:

- **3 were real**: #1690 barrier invariants, restored to do-plan-critique verbatim
- **The rest were test-side**: retargeted to assert the same contract at its new seam location
- Several were **already failing against origin/main**: PR #1806 had genericized a skill without updating them

Result: full 9-file content-test set 141 passed; 11 neighboring files 244 passed. The 76 remaining `tests/unit/` failures are all pre-existing on main and none touch skills.

Dispositions for #1883: placement and seams

Flagged, aggregated, and **deliberately not executed**. These need human sign-off:

Item	Question
do-sdlc	Doc says Bucket C project-only; it lives in skills-global. Which is right?
do-discover-paths rule-13 gap	Extend the coupling-signal set to catch <code>python tools/...</code> ?
analyze	Ship the strategic-analyst agent definition, or keep inline fallback only?
reclassify	Seam probe landed in pass 2; confirm it stays global
Model tiers	build-agent → fable; authenticity-pass → sonnet

Dispositions for #1883: merges and orphans

Merge candidates

- audit-models → do-oop-audit
- do-design-audit vs do-design-review: pick a survivor
- new-audit-skill → fold into new-skill
- claude-standards: carve the skills domain out, delegate to do-skills-audit
- audit-tools vs audit-next-tool: one owner

User-level orphans

No repo source, need a decision:

- audit-next-tool
- do-design-review
- get-telegram-messages
- searching-message-history
- sentry-cli
- the diverged user-level sentry copy

What the fleet taught us

- **Rot is the default state.** Three skills described systems that no longer exist (checking-system-logs, prime, pthread). Nobody noticed because nothing lints prose against reality
- **Duplication taxes every invocation.** do-build paid 2x tokens per call for months; rules stated 3 times cost 3 times, forever
- **The seam convention works.** Repo specifics move to `.claude/skill-context/` and `docs/sdlc/`; the same skill body now reads cleanly from a Rust repo
- **Verification is a deliverable.** telegram and officecli gained accuracy by running every command; five skills earned "untouched" as a verdict, not a default
- **Renovations decay too.** Pass 1's reflections fix was invalidated by #1828 before the branch merged

Summary

1. **All 60 skills renovated in two passes:** 7 monoliths and 65 content warnings went to zero, and the description budget dropped 36 percent while keeping every trigger phrasing
2. **Pass 2 changed what the skills say, never what they do:** goal-first bodies, deduplicated rules, verified claims, with commands, schemas, gates, and probe sentences kept byte-exact
3. **The judgment calls are queued, not buried:** every merge candidate, placement question, and orphan is aggregated in #1883 for human sign-off

What happens next

- **Human review before merge.** Description and trigger-surface changes affect skill firing on every machine
- After merge, the next `/update` **re-syncs hardlinks**, clearing the 47 transient rule-20 divergence warnings
- **#1883 disposition pass**: merges, moves, orphans, and model tiers, each needing a yes or no

Read more: PR #1894 · issue #1883 · `do-skills-audit/references/rubric.md` · `docs/features/skill-context-convention.md`